

Self organizing systems WS14 - Exercise 1

Prisoners dilemma

Dragan Avramovski
e1426093@student.tuwien.ac.at

Richard Plangger
e1025637@student.tuwien.ac.at

	P_1 accuses	P_1 stays silent
P_2 accuses	$P_1 : 2, P_2 : 2$	$P_1 : 3, P_2 : 0$
P_2 stays silent	$P_1 : 0, P_2 : 2$	$P_1 : 1, P_2 : 1$

Table 1: Prisoners' Dilemma. Sentences in years depending on the decision of the each prisoner

ABSTRACT

The purpose of this assignment is to show the strengths of Multi-Agent systems using the an agent platform. The Prisoners Dilemma problem is implemented with three independent agents that interact. Various agent behaviors are designed to mimic actions of a real human. A more complex is implemented using Bayes' Theorem. Some behavior combinations are used to run an iterative game. The outcomes of these rounds reveal that although the problem seems very simple, it is very hard to trick the opponent agent in the game. Our implementation does not implement a very complex communication structure, but it was a good opportunity to learn the paradigm and see how communication patterns in these system can be built.

Keywords

Self organizing systems, Agent, Agent Platform

1. PROBLEM DESCRIPTION

The Prisoners' Dilemma problem was chosen and implemented in JADE¹. In this problem, two criminals are convicted for a crime but still not sentenced as there is not enough evidence. Both criminals are interrogated individually and are aware of the consequences. The greatest challenge they face is what action to take. Depending on their decision it is judged how many years each of them should spent in jail. The possible outcome matrix of the problem is shown with Table 1 and it contains all the possible years for a prisoner of the game in respect to the chosen actions.

¹<http://jade.tilab.com/> Last accessed 11.12.2014

There are two actions available: Accuse or Stay Silent and these actions define the possible outcome matrix shown in Table 1. The possible outcomes reveal two main concepts: Rational and Irrational. Being rational means choosing the dominant strategy which is to play Accuse. It is dominant as only by playing rational a prisoner could get away with nothing. If they both play their dominant strategies the game ends up in equilibrium ($P_1=2$ and $P_2=2$). Being irrational means they always loose but it could lead them to the optimum result for both, which is to stay silent ($P_1=1$ and $P_2=1$). This strategy is a high risk for both of them as they could end up with maximum sentence. There are three agents implemented: Judge and two prisoners. The Judge agent is a very basic one and is implemented with a static behavior. The judge announces the sentence to each prisoner individually after both have told them their action. The behavior of the prisoner agents is implemented with different strategies:

1. Random. The prisoner decides using a random generator.
2. Static. The prisoner decides his action depending on his starting value.
3. Titfortat. This behavior only considered the last decision of the judge. It replays the action of the other prisoner. Initially the prisoner plays silent to start the first round.
4. Aggressive. This behavior tries to trick the other prisoner. Initially this behavior plays accuse. It tracks a window of the last x rounds and sums up the years it got for it. Having this value below a certain threshold y it stays aggressive. Depending on the initial configuration it plays another behavior if the threshold is exceeded.
5. Optimistic. The prisoner forgives and has a tendency to cooperate. If it finds out that his optimistic behavior is abused, he changes into aggressive mode.
6. Bayes. A learning Agent implemented with conditional probability - Bayes' Theorem [1]. The agent trains a model using previous rounds and plays a certain strategy. After his model has been trained he computes the probabilities and tries to take advantage of his knowledge.

P_1 decisions	P_2 decisions
<i>Accuse</i>	<i>Silent</i>
<i>Accuse</i>	<i>Silent</i>
<i>Accuse</i>	<i>Accuse</i>
<i>Accuse</i>	<i>Accuse</i>
<i>Accuse</i>	<i>Accuse</i>
<i>Accuse</i>	<i>Accuse</i>
<i>Silent</i>	<i>Accuse</i>
<i>Silent</i>	<i>Silent</i>
<i>Silent</i>	<i>Silent</i>

Table 2: Sample training set

	P_1 accuses	P_1 stays silent
P_2 accuses	4	2
P_2 stays silent	1	2

Table 3: Frequency Values based on the Training set in Table 2

2. LEARNING AGENT - BAYES

As we put the most effort into this behavior, this section is dedicated to the “learning” agent. In the following examples P_2 learns and tries to predict the next actions of P_1 . The requirement is implemented by using conditional Probability - Bayes Rule.

$$P(C|x) = P(x|C) \frac{P(C)}{P(x)}$$

C is the class/action to be predicted. Either *Accuse* or *Silent* for P_1 . x is the rounds the algorithm has seen so far. *Accuse* or *Silent* from P_2 .

2.1 Example:

Q: What is the probability of P_1 playing *Silent* given that P_2 played *Accuse*?

$$P(P_1:Silent|P_2:Accuse)=?$$

The class $C = Silent$ given the predictor data $x = Accuse$. This question can only be answered using prior knowledge given in the data set in Table 2.

To extract frequency and likelihood information the knowledge from the last rounds is counted and summed up in Table 3.

Then by following the Bayes algorithm the likely hood must be calculated. The results are shown in Table 4.

Q: What is the probability of P_1 playing *Silent* given that P_2 played *Accuse*?

$$P(P_1:Silent|P_2:Accuse)=0.3333$$

Q: What is the probability of P_1 playing *Accuse* given that P_2 played *Accuse*?

$$P(P_1:Accuse|P_2:Accuse)=0.6666$$

3. IMPLEMENTATION AND SETUP

Our implementation uses the JADE agent framework. We use gradle as a build and dependency management tool, but we also provide a simple packaged jar. There are three actors involved on an actor platform. The judge J is the communication partner for both prisoners (P_1, P_2). Both prisoners are not allowed to communicate and only send their messages to the judge.

The behavior of J is static and implements the logic we explained earlier. It can be found in the class `JudgeBehaviour`. It is cyclic (never ending until manual intervention) and awaits two messages from different prisoners and sends back the sentence.

```
$ java -cp libs/jade.jar:libs/dilemma.jar \
    jade.Boot -name dilemma -gui \
    -agents judge:sos.agent.Judge
```

The previous command starts the agent and registers the judge Agent to the dilemma platform. J is the central role, thus it provides the gui to observe the platform.

The prisoners are started in a similar way:

```
$ java -cp libs/jade.jar:libs/dilemma.jar \
    jade.Boot -container -agents \
    p1:sos.agent.Prisoner(static,accuse) \
    -Drounds=20
```

This starts a prisoner named `p1` and plays a static behavior sending *accuse* for 20 rounds. A full list of all possible combinations can be retrieved by removing the parameters for the prisoner. Note that for the platform to work you need to start two prisoners with the name `p1` and `p2`. After a prisoner has completed his rounds he will print the stats of his iteration including the stats of the other prisoner.

3.1 Communication

While familiarizing with the JADE platform we learned that we basically have several options on how to let the agents communicate.

- Let the judge provide a service (register on the platforms yellow pages) and let the prisoners lookup the service and use it
- Judge and prisoners have unique names. This allows to directly address the other agents.
- Send messages into the platform and flag them with certain properties (e.g. Ontology: Prisoner Dilemma) and let all participants filter the messages.

While all of them have their strengths and weaknesses we decided to directly address the agents. The motivation of this was mostly because in real life this works quite the same. You know which court you must go to to pledge innocent and receive your sentence. This allowed us to filter the messages in the behavior using the AID. The behavior of sending their decision and receiving the sentence for every prisoner is the

	P_1 accuses	P_1 stays silent	$P(x)$
P_2 accuses	4/5	2/4	$P(P_2:Accuse)=6/9$
P_2 stays silent	1/5	2/4	$P(P_2:Silent)=3/9$
$P(C)$	$P(P_1:Accuse)=5/9$	$P(P_1:Silent)=4/9$	

Table 4: Likelihood Values based on the Training set in Table 2

same, thus we abstracted this behavior into the base class `sos.agent.behaviour.BasePrisonerBehaviour`.

The programming style of receiving messages in JADE is not very intuitive at first glance. We would have expected blocking behavior for receiving messages, but one must either poll for the message (bad performance) or change the control flow and call `block()` and return the action method. Rethinking this our conclusion is that it resembles behavior in the real world. If you send out a message and want your answer back, the answer often is asynchronous. In a restaurant whenever you order a dish, you could constantly poll the waiter if it is ready. But this does not make sense and would drive the waiter crazy. The normal behavior would be to wait until you receive a new message (your dish arrives) and then process it (eat it).

Figure 1 shows a communication round as a UML sequence diagram. One round is illustrated. The judge awaits two decisions (accuse|silent) from two different prisoners. He then later replies with the sentence.

Figure 2 show a subset of the classes that we have implemented.² Our base prisoner behavior implements the `action()` method to communicate with the judge. The selected behavior (depending on the start parameter of the agent) then gets the list of previous rounds. This is the knowledge he operates on. In the different behaviors we have used this information to try to implement rational agents, irrational agents.

4. TESTING OUR IMPLEMENTATION

After presenting your implementation we would like to examine the agents and their behaviors. In the following we will evaluate some behavior combinations running the prisoners for 100 rounds. The most interesting property is the sum of years, and the sum of years for each specific prisoner.

Here is a list of certain combination we believe that are interesting.

1. P_1 random, P_2 Titfortat. Since Titfortat replays the last round, and the other prisoner plays random, P_2 should have nearly equal years than P_1 .
2. P_1 optimistic, P_2 aggressive. The aggressive prisoner tries to gain the most out of the optimistic behavior. We suspect that the aggressive might have a slight advantage, because he can trick the optimistic prisoner for some rounds.
3. P_1 bayes, P_2 random. P_1 after having seen 20 rounds, should keep a window of 20 rounds that he uses to

²Parameter for the function `decide` is a list of rounds (written as `Round[*]`).

predict the other action. Here we also expect that bayes should have a slight advantage, but should not be worse than P_2 .

4. P_1 bayes, P_2 optimistic. Similar to the previous point, bayes should train a model, that predicts the optimistic behavior. We expect the two to minimize the years (after the initial random phase has worn out).
5. P_1 bayes, P_2 aggressive. P_1 should learn the aggressive behavior and also play aggressive until the end.
6. P_1 bayes, P_2 titfortat. We think that in this behavior it greatly depends on the initial model how bayes behaves.

Table 5 show the results of our test cases. The numbers written in subscript on each year denote the round. $(111)_2$ means 111 years on the second test run.

As we have expected for (1.) the years for both prisoners are as if both played random. They have nearly the same amount of years after 100 rounds. In (2.) we observe the same behavior. But (2.) also depends on the initial model. Since P_1 plays random for the first 20 rounds it determines the model. In round 3 random yielded silent very often which drove the model of bayes into a state where it predicted that silent is the best choice.

In the third combination (3.) the optimistic agent clearly lost. In each test run he got at least 30 years more than the aggressive one. It mostly depends on the parameter settings for both how bad/good the optimistic agent behaves. But we believe that the optimistic strategy is quite good, because he only allows the other prisoner to betray him until his trust level falls below zero. After that trust has to be gained again.

Test variation (4.) is very interesting. We expected P_1 to have a slight advantage. In fact he won every round with a mean value of 75 years less. For the random implementation we used the `SecureRandom` class to get better random numbers. We do not know the cause of this. Our best guess is that the sliding window of Bayes helps him to keep an up

³ P_1 uses 20 rounds to train and maintain his model

⁴Parameters for P_1 are trust level = 10, P_2 plays static silent when years are above the threshold. In round 1, window = 10, threshold = 0.6, round 2: window = 20, threshold = 0.6, round 3: window = 20, threshold = 0.5

⁵ P_1 uses sliding window of 10 rounds to train and maintain his model

⁶ P_2 has a trust level = 10, P_1 uses sliding window of 20 rounds to train and maintain his model

⁷Both consider only the last 20 rounds, P_2 uses a threshold of 0.6 and plays static silent if he exceeds the threshold, P_1 uses 20 rounds to train and maintain his model

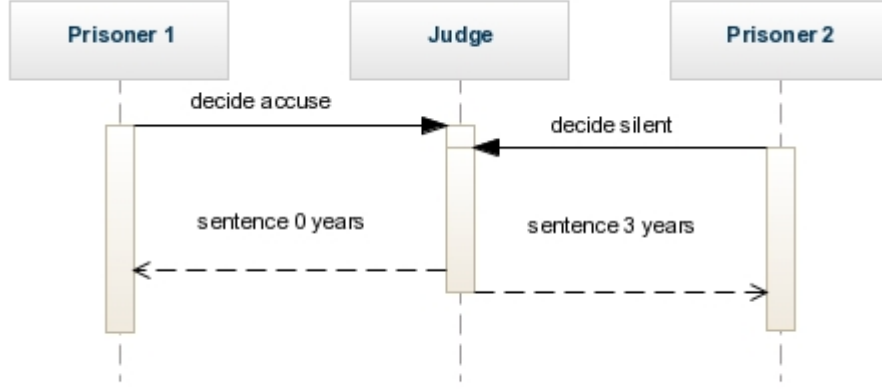


Figure 1: Sequence diagram of one round of communication

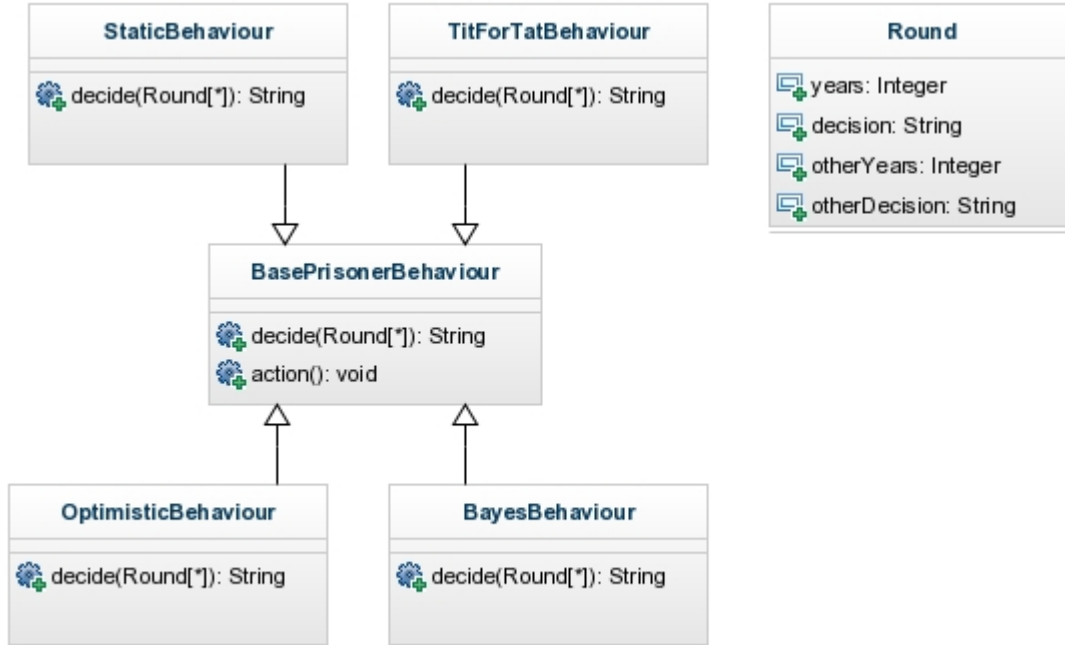


Figure 2: Subset of the classes as UML class diagram

Modus	P_1 years	P_2 years	Sum years
1. P_1 random, P_2 titfortat	(147) ₁ , (140) ₂ , (148) ₃	(147) ₁ , (143) ₂ , (142) ₃	(294) ₁ , (283) ₂ , (290) ₃
2. P_1 bayes, P_2 titfortat ³	(182) ₁ , (187) ₂ , (111) ₃	(188) ₁ , (184) ₂ , (111) ₃	(370) ₁ , (371) ₂ , (222) ₃
3. P_1 optimistic, P_2 aggressive ⁴	(177) ₁ , (182) ₂ , (183) ₃	(141) ₁ , (140) ₂ , (120) ₃	(318) ₁ , (322) ₂ , (303) ₃
4. P_1 bayes, P_2 random ⁵	(120) ₁ , (108) ₂ , (135) ₃	(204) ₁ , (222) ₂ , (162) ₃	(324) ₁ , (330) ₂ , (297) ₃
5. P_1 bayes, P_2 optimistic ⁶	(155) ₁ , (150) ₂ , (154) ₃	(203) ₁ , (207) ₂ , (205) ₃	(358) ₁ , (357) ₂ , (359) ₃
6. P_1 bayes, P_2 aggressive ⁷	(194) ₁ , (211) ₂ , (198) ₃	(140) ₁ , (142) ₂ , (150) ₃	(334) ₁ , (353) ₂ , (348) ₃

Table 5: Testing the different combinations using our system. Iteration count is always 100. Bayes behavior uses random decisions until he is has enough data (20 rounds)

to date model and that the random class is not that random at all.

Similar to (3.), in test variation (5.) the optimistic agent loses. Here he even gains more years (around 50) than in (3.) The prediction is really good and keeps P_1 from gaining years when the trust level of P_2 drops below zero.

The last test variation (6.) is the only that has beaten the Bayes prediction. An aggressive agent P_2 disguises himself as silent prisoner after he has earned played aggressive. In the second phase P_1 predicts P_2 to stay silent (he did so in the last rounds) P_2 will drop into aggressive behavior again.

5. CONCLUSION

In this assignment we have implemented a simple agent system where two prisoners are able to cooperate and/or compete against each other. The JADE interface to implement an agent is simple and stable. It would be interesting to evaluate the platform in more detail and see more complex communication patterns that are possibly split amongst sever computing nodes. In addition to that we have implemented several behaviors and let them compete against each other. Some results surprised us. It seems to be a very easily understandable problem, but in our opinion it is not so easy to trick other agents. The self organizing aspect of each agent is also very interesting. Because the interface is very simple you need to communicate to acquire more information. An agent alone can do very little things, but his capabilities increase tremendously when he gets new information.

6. REFERENCES

- [1] GK Gupta. *Introduction to data mining with case studies*. PHI Learning Pvt. Ltd., 2011.